

A Guided Tour of Chapter 11: Batch RL: Experience Replay, DQN, LSPI

Ashwin Rao

ICME, Stanford University

Incremental RL makes inefficient use of training data

- Incremental versus Batch RL in the context of fixed finite data
- Let's understand the difference for the simple case of MC Prediction
- Given fixed finite sequence of trace experiences yielding training data:

$$\mathcal{D} = [(S_i, G_i) | 1 \leq i \leq n]$$

- Incremental MC estimates $V(s; \mathbf{w})$ using $\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{w})$ for each data pair:

$$\mathcal{L}_{(S_i, G_i)}(\mathbf{w}) = \frac{1}{2} \cdot (V(S_i; \mathbf{w}) - G_i)^2$$

$$\nabla_{\mathbf{w}} \mathcal{L}_{(S_i, G_i)}(\mathbf{w}) = (V(S_i; \mathbf{w}) - G_i) \cdot \nabla_{\mathbf{w}} V(S_i; \mathbf{w})$$

$$\Delta \mathbf{w} = \alpha \cdot (G_i - V(S_i; \mathbf{w})) \cdot \nabla_{\mathbf{w}} V(S_i; \mathbf{w})$$

- n updates are performed in sequence for $i = 1, 2, \dots, n$
- Uses update method of FunctionApprox for each data pair (S_i, G_i)
- Incremental RL makes inefficient use of available training data \mathcal{D}
- Essentially each data point is “discarded” after being used for update

Batch MC Prediction makes efficient use of training data

- Instead we'd like to estimate the Value Function $V(s; \mathbf{w}^*)$ such that

$$\begin{aligned}\mathbf{w}^* &= \arg \min_{\mathbf{w}} \frac{1}{n} \cdot \sum_{i=1}^n \frac{1}{2} \cdot (V(S_i; \mathbf{w}) - G_i)^2 \\ &= \arg \min_{\mathbf{w}} \mathbb{E}_{(S,G) \sim \mathcal{D}} \left[\frac{1}{2} \cdot (V(S; \mathbf{w}) - G)^2 \right]\end{aligned}$$

- This is the solve method of FunctionApprox on training data \mathcal{D}
- This approach to RL is known as *Batch RL*
- solve by doing updates with repeated use of available data pairs
- Each update using random data pair $(S, G) \sim \mathcal{D}$

$$\Delta \mathbf{w} = \alpha \cdot (G - V(S; \mathbf{w})) \cdot \nabla_{\mathbf{w}} V(S; \mathbf{w})$$

- This will ultimately converge to desired value function $V(s; \mathbf{w}^*)$
- Repeated use of available data known as *Experience Replay*
- This makes more efficient use of available training data \mathcal{D}

Batch TD Prediction makes efficient use of *Experience*

- In Batch TD Prediction, we have experience \mathcal{D} available as:

$$\mathcal{D} = [(S_i, R_i, S'_i) | 1 \leq i \leq n]$$

- Where (R_i, S'_i) is the pair of reward and next state from a state S_i
- So, Experience \mathcal{D} in the form of finite number of atomic experiences
- This is represented in code as an Iterable [TransitionStep [S]]
- Parameters updated with repeated use of these atomic experiences
- Each update using random data pair $(S, R, S') \sim \mathcal{D}$

$$\Delta \mathbf{w} = \alpha \cdot (R + \gamma \cdot V(S'; \mathbf{w}) - V(S; \mathbf{w})) \cdot \nabla_{\mathbf{w}} V(S; \mathbf{w})$$

- This is TD Prediction with Experience Replay on Finite Experience \mathcal{D}

Batch TD(λ) Prediction

- In Batch TD(λ) Prediction, given finite number of trace experiences

$$\mathcal{D} = [(S_{i,0}, R_{i,1}, S_{i,1}, R_{i,2}, S_{i,2}, \dots, R_{i,T_i}, S_{i,T_i}) | 1 \leq i \leq n]$$

- Parameters updated with repeated use of these trace experiences
- Randomly pick trace experience (say indexed i) $\sim \mathcal{D}$
- For trace experience i , parameters updated at each time step t :

$$\mathbf{E}_t = \gamma \lambda \cdot \mathbf{E}_{t-1} + \nabla_{\mathbf{w}} V(S_{i,t}; \mathbf{w})$$

$$\Delta \mathbf{w} = \alpha \cdot (R_{i,t+1} + \gamma \cdot V(S_{i,t+1}; \mathbf{w}) - V(S_{i,t}; \mathbf{w})) \cdot \mathbf{E}_t$$

The Deep Q-Networks (DQN) Control Algorithm

DQN uses **Experience Replay** and **fixed Q-learning targets**.

At each time t for each episode:

- Given state S_t , take action A_t according to ϵ -greedy policy extracted from Q-network values $Q(S_t, a; \mathbf{w})$
- Given state S_t and action A_t , obtain reward R_{t+1} and next state S_{t+1}
- Store atomic experience $(S_t, A_t, R_{t+1}, S_{t+1})$ in replay memory \mathcal{D}
- Sample random mini-batch of atomic experiences $(s_i, a_i, r_i, s'_i) \sim \mathcal{D}$
- Update Q-network parameters \mathbf{w} using Q-learning targets based on “frozen” parameters \mathbf{w}^- of *target network*

$$\Delta \mathbf{w} = \alpha \cdot \sum_i (r_i + \gamma \cdot \max_{a'_i} Q(s'_i, a'_i; \mathbf{w}^-) - Q(s_i, a_i; \mathbf{w})) \cdot \nabla_{\mathbf{w}} Q(s_i, a_i; \mathbf{w})$$

- $S_t \leftarrow S_{t+1}$

Parameters \mathbf{w}^- of target network infrequently updated to values of Q-network parameters \mathbf{w} (hence, Q-learning targets treated as “frozen”)

Least-Squares RL Prediction

- Batch RL Prediction for general function approximation is iterative
- Uses experience replay and gradient descent
- We can solve directly (without gradient) for linear function approx
- Define a sequence of feature functions $\phi_j : \mathcal{X} \rightarrow \mathbb{R}, j = 1, 2, \dots, m$
- Parameters w is a weights vector $\mathbf{w} = (w_1, w_2, \dots, w_m) \in \mathbb{R}^m$
- Value Function is approximated as:

$$V(s; \mathbf{w}) = \sum_{j=1}^m \phi_j(s) \cdot w_j = \phi(s) \cdot \mathbf{w}$$

where $\phi(s) \in \mathbb{R}^m$ is the feature vector for state s

Least-Squares Monte-Carlo (LSMC)

- Loss function for Batch MC Prediction with data $[(S_i, G_i) | 1 \leq i \leq n]$:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2n} \cdot \sum_{i=1}^n \left(\sum_{j=1}^m \phi_j(S_i) \cdot w_j - G_i \right)^2 = \frac{1}{2n} \cdot \sum_{i=1}^n (\phi(S_i) \cdot \mathbf{w} - G_i)^2$$

- The gradient of this Loss function is set to 0 to solve for \mathbf{w}^*

$$\sum_{i=1}^n \phi(S_i) \cdot (\phi(S_i) \cdot \mathbf{w}^* - G_i) = 0$$

- \mathbf{w}^* is solved as $\mathbf{A}^{-1} \cdot \mathbf{b}$
- $m \times m$ Matrix \mathbf{A} is accumulated at each data pair (S_i, G_i) as:

$$\mathbf{A} \leftarrow \mathbf{A} + \phi(S_i) \otimes \phi(S_i) \text{ (note: } \otimes \text{ means Outer-Product)}$$

- m -Vector \mathbf{b} is accumulated at each data pair (S_i, G_i) as:

$$\mathbf{b} \leftarrow \mathbf{b} + \phi(S_i) \cdot G_i$$

- Sherman-Morrison incremental inverse can be done in $O(m^2)$

Least-Squares Temporal-Difference (LSTD)

- Loss func for Batch TD Prediction with data $[(s_i, r_i, s'_i) | 1 \leq i \leq n]$:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2n} \cdot \sum_{i=1}^n (\phi(s_i) \cdot \mathbf{w} - (r_i + \gamma \cdot \phi(s'_i) \cdot \mathbf{w}))^2$$

- The semi-gradient of this Loss function is set to 0 to solve for \mathbf{w}^*

$$\sum_{i=1}^n \phi(s_i) \cdot (\phi(s_i) \cdot \mathbf{w}^* - (r_i + \gamma \cdot \phi(s'_i) \cdot \mathbf{w}^*)) = 0$$

- \mathbf{w}^* is solved as $\mathbf{A}^{-1} \cdot \mathbf{b}$
- $m \times m$ Matrix \mathbf{A} is accumulated at each atomic experience (s_i, r_i, s'_i) :
 $\mathbf{A} \leftarrow \mathbf{A} + \phi(s_i) \otimes (\phi(s_i) - \gamma \cdot \phi(s'_i))$ (note: \otimes means Outer-Product)
- m -Vector \mathbf{b} is accumulated at each atomic experience (s_i, r_i, s'_i) :

$$\mathbf{b} \leftarrow \mathbf{b} + \phi(s_i) \cdot r_i$$

- Sherman-Morrison incremental inverse can be done in $O(m^2)$

LSTD(λ)

- Likewise, we can do LSTD(λ) using Eligibility Traces
- Denote the Eligibility Trace of atomic experience i as \mathbf{E}_i
- Note: \mathbf{E}_i accumulates $\nabla_{\mathbf{w}} V(s; \mathbf{w}) = \phi(s)$ in each trace experience
- When accumulating, previous step's eligibility trace discounted by $\lambda\gamma$

$$\sum_{i=1}^n \mathbf{E}_i \cdot (\phi(s_i) \cdot \mathbf{w}^* - (r_i + \gamma \cdot \phi(s'_i) \cdot \mathbf{w}^*)) = 0$$

- \mathbf{w}^* is solved as $\mathbf{A}^{-1} \cdot \mathbf{b}$
- $m \times m$ Matrix \mathbf{A} is accumulated at each atomic experience i :

$$\mathbf{A} \leftarrow \mathbf{A} + \mathbf{E}_i \otimes (\phi(s_i) - \gamma \cdot \phi(s'_i)) \text{ (note: } \otimes \text{ means Outer-Product)}$$

- m -Vector \mathbf{b} is accumulated at each atomic experience (s_i, r_i, s'_i) as:

$$\mathbf{b} \leftarrow \mathbf{b} + \mathbf{E}_i \cdot r_i$$

- Sherman-Morrison incremental inverse can be done in $O(m^2)$

Convergence of Least Squares Prediction Algorithms

On/Off Policy	Algorithm	Tabular	Linear	Non-Linear
On-Policy	MC	✓	✓	✓
	LSMC	✓	✓	-
	TD	✓	✓	✗
	LSTD	✓	✓	-
Off-Policy	MC	✓	✓	✓
	LSMC	✓	✗	-
	TD	✓	✗	✗
	LSTD	✓	✗	-

Least Squares RL Control

- To perform Least Squares RL Control, we do GPI with:
 - Policy Evaluation as Least-Squares Q-Value Prediction
 - Greedy Policy Improvement
- For MC or On-Policy TD Control, Q-Value Prediction (for policy π):

$$Q^\pi(s, a) \approx Q(s, a; \mathbf{w}^*) = \phi(s, a) \cdot \mathbf{w}^*$$

- Direct solve for \mathbf{w}^* using experience data generated using policy π
- We are interested in Off-Policy Control with Least-Squares TD
- Using the same idea as Q-Learning and with Experience Replay
- This technique is known as Least Squares Policy Iteration (LSPI)

Least Squares Policy Iteration (LSPI)

- In each iteration of GPI, we operate with a function approximation

$$Q(s, a; \mathbf{w}) = \phi(s, a) \cdot \mathbf{w} = \sum_{j=1}^m \phi_j(s, a) \cdot w_j$$

- Deterministic policy π_D (*target policy* for this iteration) is given by:

$$\pi_D(s) = \arg \max_a Q(s, a; \mathbf{w})$$

- Sample mini-batch of experiences (s_i, a_i, r_i, s'_i) from replay memory \mathcal{D}
- Goal of the iteration is to solve for weights \mathbf{w}' to minimize:

$$\begin{aligned} \mathcal{L}(\mathbf{w}') &= \sum_i (Q(s_i, a_i; \mathbf{w}') - (r_i + \gamma \cdot Q(s'_i, \pi_D(s'_i); \mathbf{w}')))^2 \\ &= \sum_i (\phi(s_i, a_i) \cdot \mathbf{w}' - (r_i + \gamma \cdot \phi(s'_i, \pi_D(s'_i)) \cdot \mathbf{w}'))^2 \end{aligned}$$

- \mathbf{w}' would give the parameters of the next iteration's function approx

Least Squares Policy Iteration (LSPI)

- We set the semi-gradient of $\mathcal{L}(\mathbf{w}')$ to 0

$$\sum_i \phi(s_i, a_i) \cdot (\phi(s_i, a_i) \cdot \mathbf{w}' - (r_i + \gamma \cdot \phi(s'_i, \pi(s'_i)) \cdot \mathbf{w}')) = 0$$

- \mathbf{w}' is solved as $\mathbf{A}^{-1} \cdot \mathbf{b}$
- $m \times m$ Matrix \mathbf{A} is accumulated at each experience (s_i, a_i, r_i, s'_i) :

$$\mathbf{A} \leftarrow \mathbf{A} + \phi(s_i, a_i) \otimes (\phi(s_i, a_i) - \gamma \cdot \phi(s'_i, \pi_D(s'_i)))$$

- m -Vector \mathbf{b} is accumulated at each experience (s_i, a_i, r_i, s'_i) as:

$$\mathbf{b} \leftarrow \mathbf{b} + \phi(s_i, a_i) \cdot r_i$$

- Shermann-Morrison incremental inverse can be done in $O(m^2)$
- This least-squares solution of \mathbf{w}' (Prediction) is known as *LSTDQ*
- GPI with LSTDQ and greedy policy improvement known as *LSPI*

Convergence of Control Algorithms

Algorithm	Tabular	Linear	Non-Linear
MC Control	✓	(✓)	✗
SARSA	✓	(✓)	✗
Q-Learning	✓	✗	✗
LSPI	✓	(✓)	-

(✓) means it chatters around near-optimal Value Function

LSPI for Optimal Exercise of American Options

- American Option Pricing is Optimal Stopping, and hence an MDP
- So can be tackled with Dynamic Programming or RL algorithms
- But let us first review the mainstream approaches
- For some American options, just price the European, eg: vanilla call
- When payoff is not path-dependent and state dimension is not large, we can do backward induction on a binomial/trinomial tree/grid
- Otherwise, the standard approach is [Longstaff-Schwartz algorithm](#)
- Longstaff-Schwartz algorithm combines 3 ideas:
 - Valuation based on Monte-Carlo simulation
 - Function approximation of continuation value for in-the-money states
 - Backward-recursive determination of early exercise states
- We consider LSPI as an alternative approach for American Pricing

Ingredients of Longstaff-Schwartz Algorithm

- m Monte-Carlo paths indexed $i = 0, 1, \dots, m - 1$
- $n + 1$ time steps indexed $j = n, n - 1, \dots, 1, 0$ (we move back in time)
- Infinitesimal Risk-free rate at time t_j denoted r_{t_j}
- Simulation paths (based on risk-neutral process) of underlying security prices as input 2-dim array $SP[i, j]$
- At each time step, $CF[i]$ is PV of current+future cashflow for path i
- $s_{i,j}$ denotes state for $(i, j) := (\text{time } t_j, \text{price history } SP[i, : (j + 1)])$
- $Payoff(s_{i,j})$ denotes Option Payoff at (i, j)
- $\phi_0(s_{i,j}), \dots, \phi_{r-1}(s_{i,j})$ represent feature functions (of state $s_{i,j}$)
- w_0, \dots, w_{r-1} are the regression weights
- Regression function $f(s_{i,j}) = w^T \cdot \phi(s_{i,j}) = \sum_{l=0}^{r-1} w_l \cdot \phi_l(s_{i,j})$
- $f(\cdot)$ is estimate of continuation value for in-the-money states

The Longstaff-Schwartz Algorithm

Algorithm 0.1: LONGSTAFFSCHWARTZ($SP[0 : m, 0 : n + 1]$)

comment: $s_{i,j}$ is shorthand for state at $(i,j) := (t_j, SP[i, : (j + 1)])$

$CF[0 : m] \leftarrow [Payoff(s_{i,n}) \text{ for } i \text{ in range}(m)]$

for $j \leftarrow n - 1$ **downto** 1

do $\left\{ \begin{array}{l} CF[0 : m] \leftarrow CF[0 : m] \cdot e^{-r_{t_j}(t_{j+1}-t_j)} \\ X \leftarrow [\phi(s_{i,j}) \text{ for } i \text{ in range}(m) \text{ if } Payoff(s_{i,j}) > 0] \\ Y \leftarrow [CF[i] \text{ for } i \text{ in range}(m) \text{ if } Payoff(s_{i,j}) > 0] \\ w \leftarrow (X^T \cdot X)^{-1} \cdot X^T \cdot Y \\ \textbf{comment:} \text{ Above regression gives estimate of continuation value} \\ \textbf{for } i \leftarrow 0 \textbf{ to } m - 1 \\ \quad \textbf{do } CF[i] \leftarrow Payoff(s_{i,j}) \textbf{ if } Payoff(s_{i,j}) > w^T \cdot \phi(s_{i,j}) \end{array} \right.$

$exercise \leftarrow Payoff(s_{0,0})$

$continue \leftarrow e^{-r_0(t_1-t_0)} \cdot mean(CF[0 : m])$

return ($\max(exercise, continue)$)

RL as an alternative to Longstaff-Schwartz

- RL is straightforward if we clearly define the MDP
- *State* is [Current Time, History of Underlying Security Prices]
- *Action* is Boolean: Exercise (i.e., Stop) or Continue
- *Reward* always 0, except upon Exercise (= Payoff)
- *State*-transitions based on Underlying Security's Risk-Neutral Process
- Key is function approximation of state-conditioned continuation value
- Continuation Value \Rightarrow Optimal Stopping \Rightarrow Option Price
- We outline two RL Algorithms:
 - Least Squares Policy Iteration (LSPI)
 - Fitted Q-Iteration (FQI)
- Both Algorithms are batch methods solving a linear system
- Reference: [Li, Szepesvari, Schuurmans paper](#)

LSPI customized for American Options Pricing

- a is e (exercise) or c (continue), s is $s_{i,j}$, s' is $s_{i,j+1}$
- r is $\gamma \cdot \text{Payoff}(s_{i,j+1})$ if $\pi(s_{i,j+1}) = e$ and $r = 0$ if $\pi(s_{i,j+1}) = c$
- We set $Q(s_{i,j}, e) = \text{Payoff}(s_{i,j})$ (not to be learnt)
- We set $Q(s_{i,j}, c; w) = w^T \cdot \phi(s_{i,j})$ (to be learnt)
- This requires us to set: $x(s_{i,j}, c) = \phi(s_{i,j})$ and $x(s_{i,j}, e) = 0$
- When $\pi(s_{i,j+1}) = c$, i.e., when $w^T \cdot \phi(s_{i,j+1}) \geq \text{Payoff}(s_{i,j+1})$
 - A update is: $\phi(s_{i,j}) \cdot (\phi(s_{i,j}) - \gamma \cdot \phi(s_{i,j+1}))^T$
 - B update is: 0
- When $\pi(s_{i,j+1}) = e$, i.e., when $w^T \cdot \phi(s_{i,j+1}) < \text{Payoff}(s_{i,j+1})$
 - A update is: $\phi(s_{i,j}) \cdot (\phi(s_{i,j}) - \gamma \cdot 0)^T$
 - B update is: $\gamma \cdot \text{Payoff}(s_{i,j+1}) \cdot \phi(s_{i,j})$

LSPI for American Options Pricing

Algorithm 0.2: LSPI-AMERICANPRICING($SP[0 : m, 0 : n + 1]$)

comment: $s_{i,j}$ is shorthand for state at $(i, j) := (t_j, SP[i, : (j + 1)])$

comment: A is an $r \times r$ matrix, b and w are r -length vectors

comment: $A_{i,j} \leftarrow \phi(s_{i,j}) \cdot (\phi(s_{i,j}) - \gamma \cdot \mathbb{I}_{w^T \cdot \phi(s_{i,j+1}) \geq \text{Payoff}(s_{i,j+1})} \cdot \phi(s_{i,j+1}))^T$

comment: $b_{i,j} \leftarrow \gamma \cdot \mathbb{I}_{w^T \cdot \phi(s_{i,j+1}) < \text{Payoff}(s_{i,j+1})} \cdot \text{Payoff}(s_{i,j+1}) \cdot \phi(s_{i,j})$

$A \leftarrow 0, B \leftarrow 0, w \leftarrow 0$

for $i \leftarrow 0$ **to** $m - 1$

do $\left\{ \begin{array}{l} \text{for } j \leftarrow 0 \text{ to } n - 1 \\ \text{do } \left\{ \begin{array}{l} Q \leftarrow \text{Payoff}(s_{i,j+1}) \\ P \leftarrow \phi(s_{i,j+1}) \text{ if } j < n - 1 \text{ and } Q \leq w^T \cdot \phi(s_{i,j+1}) \text{ else } 0 \\ R \leftarrow Q \text{ if } Q > w^T \cdot P \text{ else } 0 \\ A \leftarrow A + \phi(s_{i,j}) \cdot (\phi(s_{i,j}) - e^{-r_{t_j}(t_{j+1}-t_j)} \cdot P)^T \\ B \leftarrow B + e^{-r_{t_j}(t_{j+1}-t_j)} \cdot R \cdot \phi(s_{i,j}) \end{array} \right. \\ w \leftarrow A^{-1} \cdot b, A \leftarrow 0, b \leftarrow 0 \text{ if } (i + 1) \% \text{BatchSize} == 0 \end{array} \right.$

Fitted Q-Iteration for American Options Pricing

Algorithm 0.3: FQI-AMERICANPRICING($SP[0 : m, 0 : n + 1]$)

comment: $s_{i,j}$ is shorthand for state at $(i, j) := (t_j, SP[i, : (j + 1)])$

comment: A is an $r \times r$ matrix, b and w are r -length vectors

comment: $A_{i,j} \leftarrow \phi(s_{i,j}) \cdot \phi(s_{i,j})^T$

comment: $b_{i,j} \leftarrow \gamma \cdot \max(\text{Payoff}(s_{i,j+1}), w^T \cdot \phi(s_{i,j+1})) \cdot \phi(s_{i,j})$

$A \leftarrow 0, B \leftarrow 0, w \leftarrow 0$

for $i \leftarrow 0$ **to** $m - 1$

do $\left\{ \begin{array}{l} \text{for } j \leftarrow 0 \text{ to } n - 1 \\ \text{do } \left\{ \begin{array}{l} Q \leftarrow \text{Payoff}(s_{i,j+1}) \\ P \leftarrow \phi(s_{i,j+1}) \text{ if } j < n - 1 \text{ else } 0 \\ A \leftarrow A + \phi(s_{i,j}) \cdot \phi(s_{i,j})^T \\ B \leftarrow B + e^{-r_{t_j}(t_{j+1}-t_j)} \cdot \max(\text{Payoff}(s_{i,j+1}), w^T \cdot P) \cdot \phi(s_{i,j}) \end{array} \right. \\ w \leftarrow A^{-1} \cdot b, A \leftarrow 0, b \leftarrow 0 \text{ if } (i + 1) \% \text{BatchSize} == 0 \end{array} \right.$

Feature functions

- Li, Szepesvari, Schuurmans recommend Laguerre polynomials (first 3)
- Over $S' = S_t/K$ where S_t is underlying price and K is strike
- $\phi_0(S_t) = 1, \phi_1(S_t) = e^{-\frac{S'}{2}}, \phi_2(S_t) = e^{-\frac{S'}{2}} \cdot (1 - S'), \phi_3(S_t) = e^{-\frac{S'}{2}} \cdot (1 - 2S' + S'^2/2)$
- They used these for Longstaff-Schwartz as well as for LSPI and FQI
- For LSPI and FQI, we also need feature functions for time
- They recommend
 $\phi_0^t(t) = \sin(\frac{\pi(T-t)}{2T}), \phi_1^t(t) = \log(T - t), \phi_2^t(t) = (\frac{t}{T})^2$

Key Takeaways from this Chapter

- Batch RL makes efficient use of data
- DQN uses experience replay and fixed Q-learning targets, avoiding the pitfalls of time-correlation and semi-gradient
- LSTD is a direct (gradient-free) solution of Batch TD Prediction
- LSPI is an off-policy, experience-replay Control Algorithm using LSTDQ for Policy Evaluation
- Optimal Exercise of American Options can be tackled with LSPI and Deep Q-Learning algorithms