

A Guided Tour of Chapter 10: Reinforcement Learning for Control


Ashwin Rao

ICME, Stanford University

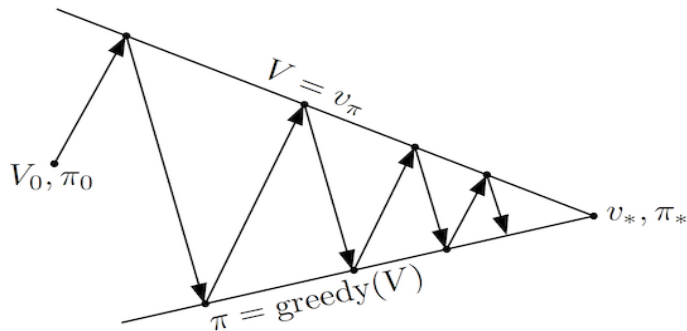
RL does not have access to a probability model

- DP/ADP assume access to probability model (knowledge of \mathcal{P}_R)
- Often in real-world, we do not have access to these probabilities
- Which means we'd need to *interact* with the *actual environment*
- Actual Environment serves up individual experiences, not probabilities
- Even if MDP model is available, model updates can be challenging
- Often real-world models end up being too large or too complex
- Sometimes estimating a *sampling model* is much more feasible
- So RL interacts with either *actual* or *simulated* environment
- Either way, we receive *individual experiences* of next state and reward
- We saw how RL Prediction learns from individual experiences
- Now we extend those ideas to RL Control: Learning Optimal VF
- We start with Tabular RL Control

Let us recall the Policy Iteration algorithm



The idea of Generalized Policy Iteration (GPI)



- Any Policy Evaluation method, Any Policy Improvement method
- Policy Evaluation estimates $V^{(\pi)}$, eg: Iterative Policy Evaluation
- Policy Improvement produces $\pi' \geq \pi$, eg: Greedy Policy Improvement
- Policy Evaluation and Policy Improvement alternate until convergence

Natural Idea: GPI with Tabular Monte-Carlo Evaluation

- Let us explore GPI with Tabular Monte-Carlo evaluation
- So we will do Policy Evaluation with Tabular MC evaluation
- And we will do the usual Greedy Policy Improvement
- But Greedy Policy Improvement requires a model of MDP

$$\pi'(s) \leftarrow \arg \max_{a \in \mathcal{A}} \{ \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{N}} \mathcal{P}(s, a, s') \cdot V^\pi(s') \}$$

- However, it works if we were working with Action-Value Function

$$\pi'(s) \leftarrow \arg \max_{a \in \mathcal{A}} Q^\pi(s, a)$$

- This means: Policy Evaluation for Action-Value Function $Q^\pi(s, a)$
- Following a policy π , update Q-value for each (S_t, A_t) each episode:

$$Count(S_t, A_t) \leftarrow Count(S_t, A_t) + 1$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{Count(S_t, A_t)} \cdot (G_t - Q(S_t, A_t))$$

ϵ -Greedy Policy Improvement

- A full Policy Evaluation with MC takes too long
- So we typically improve policy after each episode
- This can lead to some actions not being tried enough
- Which can lead to premature (greedy) domination of an action
- Which can lead to other actions getting “locked-out”
- Same as *Explore v/s Exploit* dilemma of Multi-Armed Bandit problem
- Simple solution: Perform an ϵ -Greedy Policy Improvement
- All $|\mathcal{A}|$ actions are tried with non-zero probability (for each state)
- Pick the greedy action with probability $1 - \epsilon$
- With probability ϵ , randomly choose one of the $|\mathcal{A}|$ actions

$$\text{Stochastic Policy } \pi(s, a) = \begin{cases} \frac{\epsilon}{|\mathcal{A}|} + 1 - \epsilon & \text{if } a = \arg \max_{b \in \mathcal{A}} Q(s, b) \\ \frac{\epsilon}{|\mathcal{A}|} & \text{otherwise} \end{cases}$$

ϵ -Greedy improves the policy

Theorem

For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to Q^π is an improvement, i.e., $\mathbf{V}^{\pi'}(s) \geq \mathbf{V}^\pi(s)$ for all $s \in \mathcal{N}$.

- Applying $\mathbf{B}^{\pi'}$ repeatedly (starting with \mathbf{V}^π) converges to $\mathbf{V}^{\pi'}$:

$$\lim_{i \rightarrow \infty} (\mathbf{B}^{\pi'})^i(\mathbf{V}^\pi) = \mathbf{V}^{\pi'}$$

- So the proof is complete if we prove that:

$$(\mathbf{B}^{\pi'})^{i+1}(\mathbf{V}^\pi) \geq (\mathbf{B}^{\pi'})^i(\mathbf{V}^\pi) \text{ for all } i = 0, 1, 2, \dots$$

- Increasing tower of Value Functions $[(\mathbf{B}^{\pi'})^i(\mathbf{V}^\pi) | i = 0, 1, 2, \dots]$ with repeated applications of $\mathbf{B}^{\pi'}$

Proof of ϵ -Greedy improving the policy

To prove base case (proof by induction), note: $\mathbf{B}^{\pi'}(\mathbf{V}^{\pi})(s) = Q^{\pi}(s, \pi'(s))$

$$\begin{aligned} Q^{\pi}(s, \pi'(s)) &= \sum_{a \in \mathcal{A}} \pi'(s, a) \cdot Q^{\pi}(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}|} \cdot \sum_{a \in \mathcal{A}} Q^{\pi}(s, a) + (1 - \epsilon) \cdot \max_{a \in \mathcal{A}} Q^{\pi}(s, a) \\ &\geq \frac{\epsilon}{|\mathcal{A}|} \cdot \sum_{a \in \mathcal{A}} Q^{\pi}(s, a) + (1 - \epsilon) \cdot \sum_{a \in \mathcal{A}} \frac{\pi(s, a) - \frac{\epsilon}{|\mathcal{A}|}}{1 - \epsilon} \cdot Q^{\pi}(s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(s, a) \cdot Q^{\pi}(s, a) = \mathbf{V}^{\pi}(s) \end{aligned}$$

Induction step is proved by monotonicity of \mathbf{B}^{π} operator (for any π):

Monotonicity Property of $\mathbf{B}^{\pi} : \mathbf{X} \geq \mathbf{Y} \Rightarrow \mathbf{B}^{\pi}(\mathbf{X}) \geq \mathbf{B}^{\pi}(\mathbf{Y})$

So $(\mathbf{B}^{\pi'})^{i+1}(\mathbf{V}^{\pi}) \geq (\mathbf{B}^{\pi'})^i(\mathbf{V}^{\pi}) \Rightarrow (\mathbf{B}^{\pi'})^{i+2}(\mathbf{V}^{\pi}) \geq (\mathbf{B}^{\pi'})^{i+1}(\mathbf{V}^{\pi})$



Definition

Greedy in the Limit with Infinite Exploration (GLIE):

- All state-action pairs are explored infinitely many times

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- The policy converges to a greedy policy

$$\lim_{k \rightarrow \infty} \pi_k(s, a) = \mathbb{I}_{a=\arg \max_{b \in \mathcal{A}} Q(s, b)}$$

ϵ -greedy can be made GLIE if ϵ is reduced as: $\epsilon_k = \frac{1}{k}$

GLIE Tabular Monte-Carlo Control

- Sample k -th episode using π : $\{S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T\} \sim \pi$
- For each state S_t and action A_t in the episode:

$$\text{Count}(S_t, A_t) \leftarrow \text{Count}(S_t, A_t) + 1$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{\text{Count}(S_t, A_t)} \cdot (G_t - Q(S_t, A_t))$$

- Improve policy at end of episode based on updated Q -Value function:

$$\epsilon \leftarrow \frac{1}{k}$$

$$\pi \leftarrow \epsilon\text{-greedy}(Q)$$

Theorem

GLIE Tabular Monte-Carlo Control converges to the Optimal Action-Value function: $Q(s, a) \rightarrow Q^(s, a)$.*

MC versus TD Control

- TD learning has several advantages over MC learning:
 - Lower variance
 - Online
 - Can work with incomplete traces or continuing traces
 - Generic interface of `Iterable` of atomic experiences allows for serving up atomic experiences in any order (eg: atomic experience replays)
- So use TD instead of MC in our Control loop
 - Apply TD to $Q(S, A)$ (instead of $V(S)$)
 - Use ϵ -greedy Policy Improvement
 - Update $Q(S, A)$ after each *atomic experience*
 - ϵ -greedy policy automatically updated after each atomic experience


Tabular SARSA Algorithm

- Tabular SARSA is our first TD Control algorithm
- Like Tabular MC Control, Policy Improvement is ϵ -greedy
- But here Policy Evaluation is with a TD target, as below:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \cdot (R_{t+1} + \gamma \cdot Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

- Note that $Q(S, A)$ is updated after each atomic experience
- ϵ -greedy policy automatically updated after each atomic experience
- Action A_t is chosen from State S_t based on ϵ -greedy policy
- Action A_{t+1} is chosen from State S_{t+1} based on ϵ -greedy policy
- Note: Instead of ϵ -greedy, we could employ a more sophisticated exploratory policy derived from Q -value function (ϵ -greedy is just our default simple exploratory policy derived from Q -value function)

SARSA Visualization



Convergence of Tabular SARSA

Theorem

Tabular SARSA converges to the Optimal Action-Value function, $Q(s, a) \rightarrow Q^(s, a)$, under the following conditions:*

- *GLIE sequence of policies $\pi_t(s, a)$*
- *Robbins-Monro sequence of step-sizes α_t*

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

Tabular n -step SARSA

- Tabular SARSA bootstraps the Q -Value Function with update:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

- So it's natural to extend this to bootstrapping with 2 steps ahead:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma R_{t+2} + \gamma^2 Q(S_{t+2}, A_{t+2}) - Q(S_t, A_t))$$

- Generalize to bootstrapping with $n \geq 1$ time steps ahead:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(G_{t,n} - Q(S_t, A_t))$$

- $G_{t,n}$ (call it n -step bootstrapped return) is defined as:

$$\begin{aligned} G_{t,n} &= \sum_{i=t+1}^{t+n} \gamma^{i-t-1} \cdot R_i + \gamma^n \cdot Q(S_{t+n}, A_{t+n}) \\ &= R_{t+1} + \gamma \cdot R_{t+2} + \dots + \gamma^{n-1} \cdot R_{t+n} + \gamma^n \cdot Q(S_{t+n}, A_{t+n}) \end{aligned}$$

Tabular λ -Return SARSA

- Instead of $G_{t,n}$, a valid target is a weighted-average target:

$$\sum_{n=1}^N u_n \cdot G_{t,n} + u \cdot G_t \text{ where } u + \sum_{n=1}^N u_n = 1$$

- Any of the u_n or u can be 0, as long as they all sum up to 1
- The λ -Return target is a special case of weights u_n and u

$$u_n = (1 - \lambda) \cdot \lambda^{n-1} \text{ for all } n = 1, \dots, T - t - 1$$

$$u_n = 0 \text{ for all } n \geq T - t \text{ and } u = \lambda^{T-t-1}$$

- We denote the λ -Return target as $G_t^{(\lambda)}$, defined as:

$$G_t^{(\lambda)} = (1 - \lambda) \cdot \sum_{n=1}^{T-t-1} \lambda^{n-1} \cdot G_{t,n} + \lambda^{T-t-1} \cdot G_t$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \cdot (G_t^{(\lambda)} - Q(S_t, A_t))$$

Tabular SARSA(λ)

- λ can be tuned from SARSA ($\lambda = 0$) to MC Control ($\lambda = 1$)
- Note that for $\lambda > 0$, λ -Return SARSA is an Offline Algorithm
- SARSA(λ) is online “version” of λ -Return SARSA
- Similar to TD(λ) for Prediction, SARSA(λ) uses Eligibility Traces
- Eligibility Trace for a given trace experience at time t is a function

$$E_t : \mathcal{N} \times \mathcal{A} \rightarrow \mathbb{R}_{\geq 0}$$

$$E_0(s, a) = 0, \text{ for all } s \in \mathcal{N}, a \in \mathcal{A}$$

$$E_t(s, a) = \gamma \cdot \lambda \cdot E_{t-1}(s, a) + \mathbb{I}_{S_t=s, A_t=a}, \text{ for all } s \in \mathcal{N}, a \in \mathcal{A}, \text{ for all } t$$

- Tabular SARSA(λ) performs following update at each time step t in each trace experience (for each $s \in \mathcal{N}, a \in \mathcal{A}$):

$$Q(s, a) \leftarrow Q(s, a) + \alpha \cdot (R_{t+1} + \gamma \cdot Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)) \cdot E_t(s, a)$$

Key Takeaways from this Chapter

- Bias-Variance tradeoff of TD versus MC
- MC learns the mean of the observed returns while TD learns something "deeper" - it implicitly estimates an MRP from given data and produces the Value Function of the implicitly-estimated MRP
- Understanding TD versus MC versus DP from the perspectives of:
 - "Bootstrapping"
 - "Experiencing"
- "Equivalence" of λ -Return Prediction and $TD(\lambda)$ Prediction
- TD is equivalent to $TD(0)$ and MC is "equivalent" to $TD(1)$